STELLANTIS · PRODUCT DESIGN OFFICE

Global Design AI Intake & Integration Framework

A repeatable process for proposing, validating, piloting, and scaling AI initiatives across the global Design organization — with visibility, governance, and alignment built in.

Version 1.0 · June 2026 · Owner: Global AI Design Lead (Design Functional AI Hub)

Operating principle: the Global AI Design Lead orchestrates the process. The teams closest to the problem own the work.

1. Purpose & Operating Principles

Design teams across North America, Europe, South America, and Asia-Pacific are already experimenting with AI. The risk is not too little initiative — it is fragmentation: duplicate pilots, invisible work, unscaled wins, and lessons that evaporate when a pilot ends. This framework provides one repeatable pipeline that lets any Design team move from idea to scaled capability quickly, while the organization keeps visibility, governance, and a compounding knowledge base.

This framework deliberately does not create a central team that tests every tool. It creates a lightweight operating system around the teams who already want to test, so their work is visible, comparable, governed, and reusable.

Eight operating principles

- **One front door —** every AI opportunity enters through the same global intake, regardless of region or seniority. No side doors.
- **Orchestrate, don't operate —** the Global AI Design Lead runs the process — triage, portfolio, connections, reporting. Testing, piloting, and administering solutions stays with the originating teams.
- **No ownership, no initiative —** nothing proceeds without a named Business Sponsor, Process Owner, and Pilot Owner.
- **Business problem first —** every initiative is anchored to a real workflow and a measurable benefit before any tool is discussed.
- **Search before you build —** the Design AI repository, AI Garage, and the enterprise use-case portfolio are checked at intake. Duplicates are merged, not multiplied.
- **Evidence over enthusiasm —** one scorecard, one pilot charter, one assessment standard — so a pilot in Turin is comparable to a pilot in Auburn Hills.
- **Right-sized governance —** review effort scales with risk (data sensitivity, new vendors, spend), not with organizational distance. Fast lane for low-risk work.
- **Every initiative leaves knowledge —** pilots that do not scale still produce a documented lesson. Retired is a result, not a failure.

2. Scope & Position in the Stellantis AI Operating Model

In scope

- AI tools, agents, workflows, and data products proposed by or for Design teams in NA, EE, SA, and AP/China.
- Initiatives spanning creative work (sketching, visualization, animation, CMF), design operations (reviews, research, benchmarking, program memory), and Design's interfaces to engineering, manufacturing, and marketing.
- Both new technology requests and new uses of already-approved platforms.

Out of scope

- Enterprise platform selection and operation (owned by GDPA / ICT).
- AI features of the vehicle product itself (owned by PDT).
- Personal productivity use of already-approved tools within approved environments — just use them; no intake required.

Position in the enterprise model

Design operates as one of the Global Functional AI Hubs in the Stellantis AI Transformation Office. This framework is how the Design hub executes its leadership directive — Connect, Build, Adopt, Signal, Represent — at studio level. It does not duplicate enterprise mechanisms; it plugs into them:

Enterprise mechanism	How this framework uses it
AI Garage / AI Intake Agent	Checked at intake for existing solutions; Design submissions with enterprise relevance are mirrored into the enterprise idea backlog.
GDPA platforms (Foundry, Databricks, Snowflake) & AI & Data Talent Pool	Default build path for Tier 2–3 pilots; Talent Pool engaged for build-heavy initiatives instead of hiring or shadow IT.
Data&AI Marketplace / Collibra	Design Data Products produced by initiatives are certified and published for reuse.
Stellantis AI Risk Assessment (regulatory & operational)	Triggered automatically by tier at the pilot and scale gates.
VCP / Wave	Initiatives with material recurring value are flagged in Wave under the relevant workstream (e.g., Product Cost, ER&D, Commercial & Quality).
AI & Data Ambassador network / Global AI Academy	Design's Regional AI Champions and training plans operate as part of the enterprise network (255 → 2,000 ambassadors).

3. Roles & Ownership

Every initiative must name three owners at intake. The framework stalls by design when ownership is missing — an initiative nobody owns is an initiative nobody needs.

Role	Accountable for	Notes
Business Sponsor	The business outcome: approves the value claim, unblocks resources, accepts or rejects the result.	Acts as the "AI Business Owner" for the Stellantis regulatory AI risk assessment when one is required.
Process Owner	The workflow being improved: validates the current process, the proposed future process, and owns adoption after scaling.	May be the same person as the sponsor in small teams — but the role must be explicitly accepted.
Pilot Owner / Champion	Executing the pilot: running the test, collecting measurements, reporting at checkpoints, writing the final assessment.	Owns the operational AI risk assessment with the ICT use-case lead when required. Remains in the originating team.
Global AI Design Lead	The process: intake triage, portfolio, evaluation standards, enterprise brokering, knowledge base, executive reporting.	Explicit non-responsibilities: does not test tools, does not run or administer pilots, does not own pilot results.
Regional AI Champions (NA · EE · SA · AP)	First contact for their region: help requesters complete intake, sense-check duplicates, approve Tier 1 fast-lane starts, drive local	Part of the enterprise AI & Data Ambassador network.

	adoption.	
Design Data Managers / Knowledge Stewards	Curating the repository and certifying Design Data Products per the enterprise call-to-action.	Nominated per the AI Hub onboarding requirements.
Design ICT AI Lead & Security partners	Feasibility, platform path, integration, security and compliance review at the tiers that require it.	Engaged through defined interaction points (Section 9) — not on every initiative.

RACI summary

Activity	Lead	Sponsor	Process Owner	Pilot Owner	Reg. Champion	ICT / Security
Submit intake	I	A	C	R	C	—
Duplicate & repository check	A/R	I	—	C	R	—
Evaluation & prioritization (G1)	A/R	C	C	C	C	C (Tier 2+)
Pilot execution & checkpoints	I	A	C	R	C	C (per tier)
Security / AI risk review	I	A (regulatory)	—	R	—	A (operational)
Scale decision (G4)	R (business case)	A	C	C	C	C
Knowledge capture	A	I	C	R	C	—
Executive reporting	A/R	I	I	I	C	I

R = Responsible · A = Accountable · C = Consulted · I = Informed. Scale decisions are accountable to Design leadership with the sponsor; the Lead assembles the case.

4. The Pipeline: Stages & Gates

Every initiative lives in exactly one stage. Stage status is visible to the whole Design organization through the portfolio board (Section 11).

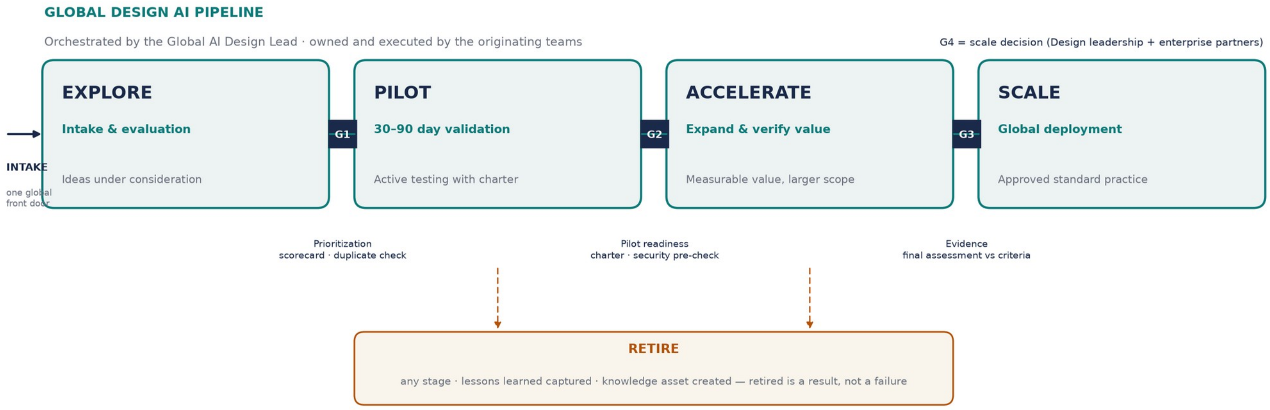

Stage	Definition	Entry requires	Exit requires
Explore	Ideas under consideration.	Complete intake with named owners.	Gate 1: scorecard disposition — proceed, redirect, merge, park, or decline.
Pilot	Active testing and validation.	Gate 2: approved pilot charter; tier-appropriate security/feasibility pre-checks done.	Gate 3: final assessment vs success criteria within 2 weeks of pilot end.
Accelerate	Initiatives demonstrating measurable value.	Pilot met success criteria; sponsor confirms value; expansion plan (more users / programs / studios).	Gate 4: scale business case — value, cost, platform path, change plan.
Scale	Approved for broader deployment.	Design leadership approval; enterprise alignment (platform, licensing, security) confirmed.	Handover to Process Owner as standard practice; adoption tracked; initiative closed into knowledge base.
Retire	Did not demonstrate sufficient value — from any stage.	Disposition decision at any gate or checkpoint.	Lessons-learned entry published; knowledge asset created. Mandatory — no silent failures.

5. Intake

One global form, designed to be completed in 15 minutes without assistance. Requesters who prefer a conversation start with their Regional AI Champion, who completes the form with them. Intake is acknowledged within 5 business days and triaged at the next monthly Design AI Council — no submission waits more than 30 days for a disposition and a named next step.

Required intake fields

Field	What good looks like
Initiative name	Short, plain language. "Benchmark image search for interiors," not "AI transformation phase 2."
Requesting organization	Studio / department submitting.

Region	NA · EE · SA · AP — or Global if inherently cross-regional.
Business problem	The pain in one paragraph, in workflow terms. No technology words required.
Current process	How the work happens today: steps, tools, time, people.
Proposed future process	How the work would happen with the solution. A sketch beats an essay.
Expected benefits	What improves — time, cost, quality, experience — and roughly how much.
Users impacted	Who and how many, at pilot and at scale.
Estimated value	Order-of-magnitude is fine at intake (hours/week, € range). Refined at evaluation.
Requested technology or platform	If known. "Unknown — need guidance" is a valid answer; the Lead brokers options.
Pilot Owner / Champion	Named individual in the requesting team.
Business Sponsor	Named individual accountable for the outcome.
Process Owner	Named individual who owns the workflow.
Data classification	Public · internal · confidential / unreleased-program · personal data. Drives the governance tier.
Existing-solution check	Confirmation the repository and AI Garage were searched; list near-matches found.
Desired timeline & reach	When the team wants to start; single studio, region, or global ambition.

The first twelve fields satisfy the framework's mandatory requirements; the final four exist to route the initiative correctly (tier, duplicates, sequencing) without a second round-trip to the requester.

6. Evaluation & Prioritization (Gate 1)

Every initiative is scored on the same ten criteria, 1–5 each, weighted as below. Scoring is done at the monthly Design AI Council by the Lead and Regional Champions, with ICT/Security consulted where flagged. The scorecard creates comparability — it does not replace judgment; the Council can override with a recorded rationale.

Criterion	Weight	What it measures
Strategic alignment	15	Advances Design Intelligence priorities and the enterprise AI strategy (data products, semantic layer, agents).
Global applicability	15	Useful beyond the originating studio; potential to become a shared capability.
Time savings	12	Hours returned to designers and studio teams.
Cost savings	10	Direct cost reduction or avoidance.
Quality improvement	10	Better design outcomes: earlier issue detection, fewer reworks, stronger reviews.
Technical feasibility	10	Achievable with available platforms

		and skills in a 30–90 day pilot.
User experience improvement	8	Reduces friction in daily design work; desirability to users.
Security considerations	8	Exposure of unreleased-program material, IP, vendor risk. Gating: a score of 1 pauses the initiative regardless of total.
Data considerations	6	Data availability, quality, classification, and rights. Gating: a score of 1 pauses the initiative.
Change management impact	6	Training burden, workflow disruption, adoption risk.

Dispositions

- **Proceed to Pilot —** score ≥ 70, owners confirmed, tier checks passed. Enters pilot queue.
- **Redirect —** an existing solution already covers the need — requester is connected to it. Logged as a win for reuse.
- **Merge —** duplicates an active initiative — teams are connected; one initiative continues with both regions.
- **Park —** promising but not now (score 50–69, missing prerequisite, or pilot capacity full). Given a revisit date — parked is not a polite no.
- **Decline —** score < 50 or gating issue with no path. Reason recorded in the repository; requester gets the rationale.

Pilot capacity is protected: as a guideline, each region runs no more than three concurrent pilots, so pilots get real attention and finish. Prioritization favors high value + low effort first; high value + high effort initiatives are sequenced with enterprise (GDPA/Talent Pool) support.

7. Pilot Standards (Gates 2–3)

Pilots are executed and owned by the originating team. The framework standardizes how pilots are framed and judged — not who runs them.

Pilot charter (required at Gate 2)

- **Objectives & hypothesis —** what we expect to improve, and why.
- **Success criteria —** quantified, agreed with the sponsor before starting; baseline measured first ("review prep takes 6 h today; target ≤ 3 h").
- **Duration —** 30–90 days; 60 is the default. Longer means the scope is too big — split it.
- **Participating users —** named participants; minimum group size agreed for the evidence to mean something.
- **Measurement approach —** how data is collected (time logs, output counts, user ratings) and by whom.
- **Checkpoint reviews —** a 15-minute mid-point report at the Design AI Council: on track / at risk / stop.
- **Final assessment —** due within two weeks of pilot end (template, Appendix D): results vs criteria, value estimate, user feedback, security/data observations, recommendation — accelerate, iterate once, or retire.

An honest "this did not work, here is why" assessment is treated as a first-class contribution to the knowledge base — it saves the next three teams from the same dead end.

8. Right-Sized Governance: Three Tiers

Governance effort scales with risk. The tier is set at intake from data classification, technology novelty, and spend — and determines which reviews are required. This is how the framework ensures appropriate review without bureaucracy.

Tier	Triggers	Required path	Target speed
Tier 1 · Fast lane	Approved tools and environments; no new data exposure; no new vendor; no incremental spend.	Regional Champion approval → log in portfolio → start. Outcome reported like any pilot.	< 1 week from intake
Tier 2 · Standard	New tool or vendor; internal (non-confidential) data; limited users; modest spend.	Full intake → scorecard → ICT feasibility check + security pre-check → pilot charter.	≤ 30 days to pilot start
Tier 3 · Elevated	Confidential / unreleased-program or personal data; new platform; integration with enterprise systems; spend above the Design-leadership threshold.	Tier 2 path + Stellantis AI Risk Assessment (regulatory before pilot where required; operational before scale) + GDPA/ICT architecture engagement + procurement.	Per enterprise review cycles — started early by the Lead

Design-specific note: sketches, models, and renders of unreleased programs are program-confidential. Any initiative moving such material outside approved environments is Tier 3 by definition — no exceptions.

9. Governance Interaction Points

Partner	Engaged when	For
ICT / Design ICT AI Lead	Tier 2+ at evaluation; all scale decisions.	Feasibility, architecture, platform path, integration, licensing.
Enterprise AI teams / GDPA	Build-heavy pilots; any initiative producing a data product; scale gate.	Platform capabilities, AI & Data Talent Pool capacity, data product publication to the Marketplace.
AI Garage	At intake (search) and after successful pilots (publish).	Discovery of existing solutions; making Design agents and use cases visible to the enterprise.
Security teams	Tier 2 pre-check; Tier 3 full review.	Vendor risk, data handling, environment approval, IP protection.
AI Governance organization	Tier 3 before pilot (regulatory) and before scale (operational).	Stellantis AI Risk Assessments; EU AI Act and regional regulatory obligations.
Regional Design leadership	Pilots affecting studio operations; all scale decisions in their region.	Resourcing, prioritization, adoption sponsorship.
VCP / Wave (with Finance)	When recurring value is material.	Flagging the initiative under the relevant VCP workstream so Design's AI value is counted.

The Lead owns these relationships and brokers every connection, so requesting teams never need to navigate the enterprise organization alone — and partner organizations get one coherent Design interlocutor instead of fifty.

10. Knowledge Management

The repository is the compounding asset of this framework: a single, searchable record of everything Design has tried, learned, and standardized. Every initiative writes to it at every gate — by the time an initiative closes, its knowledge entry already exists.

What is captured

Asset	Created at	Curated by
Use case description & intake record	Intake	Pilot Owner
Evaluation & disposition (incl. declines and redirects)	Gate 1	Lead
Pilot charter, checkpoint notes, final assessment & outcomes	Gates 2–3	Pilot Owner
Lessons learned (success and failure)	Gate 3 / Retire	Pilot Owner + Knowledge Steward
Adoption metrics for scaled solutions	Scale, ongoing	Process Owner
Training materials & recommended practices	Accelerate / Scale	Knowledge Steward + Academy
Contacts & subject-matter experts directory	Continuous	Knowledge Steward

- **Searchable and reusable —** organized by workflow, region, technology, and outcome; readable by all Design teams globally.
- **Quarterly digest —** "What Design Learned This Quarter" — distributed to all studios; the network's shared memory in practice.
- **Enterprise alignment —** the repository is itself managed as a Design Data Product — certified and registered per the enterprise data-product call-to-action, so Design's knowledge becomes a reusable enterprise asset.

11. Portfolio Management & Operating Cadence

Forum	Cadence	Participants	Decisions
Design AI Council	Monthly	Lead (chair), 4 Regional Champions, Design ICT AI Lead, rotating sponsors	Intake triage (G1), pilot starts (G2), checkpoint reviews, pilot assessments (G3), Tier escalations
Design AI Portfolio Review	Quarterly	Design leadership, Lead, Champions	Scale approvals (G4), retire decisions, investment asks, tier thresholds
Enterprise rhythms	Per AI Transformation Office calendar	Lead represents Design	Functional AI Community reporting, Steer Co inputs, VCP/Wave

			updates

The portfolio board shows every initiative by stage (Explore / Pilot / Accelerate / Scale / Retired), region, tier, and owner — visible to the entire Design organization. Visibility is the duplicate-prevention mechanism: teams can see what exists before they propose.

12. Executive Reporting

A monthly one-pager from the Council, and a quarterly executive pack aligned to the enterprise KPI framework (product usage/adoption · health & quality · incremental value · time to market).

Framework metric	Enterprise KPI category
Number of active initiatives by stage and region	Health & quality
Regional participation (initiatives and pilot owners per region)	Usage / adoption
Adoption rates of scaled solutions	Usage / adoption
Time savings achieved (validated by pilots, projected at scale)	Incremental value
Estimated business value (€ + hours), incl. Wave-flagged VCP value	Incremental value
Successful pilots (met success criteria) and pilot cycle time	Time to market
Scaled solutions in production use	Time to market / value
Knowledge assets created (use cases, lessons, training materials)	Health & quality

13. Implementation Roadmap

First 30 days	Days 31–60	Days 61–90
Publish framework + intake form. Name Regional Champions and Data Managers. Backfill all current initiatives (Vizcom, Runway, active explorations) into the portfolio — the framework launches populated, not empty. Hold Council #1: first triage.	First checkpoint reviews. Repository v1 live with backfilled lessons. Tier thresholds ratified by Design leadership. First monthly one-pager issued.	First quarterly Portfolio Review: first scale/retire decisions. First executive report into the AI Hub reporting rhythm. Publish digest #1. Confirm first Wave/VCP flags.

Appendix A — Intake Form Template

#	Field	Answer
1	Initiative name	
2	Requesting organization	
3	Region (NA / EE / SA / AP / Global)	
4	Business problem (one paragraph)	
5	Current process (steps, tools, time, people)	
6	Proposed future process	
7	Expected benefits	
8	Users impacted (pilot / at scale)	
9	Estimated value (hours, €, order of magnitude)	
10	Requested technology or platform (or "need guidance")	
11	Pilot Owner / Champion (name)	
12	Business Sponsor (name)	
13	Process Owner (name)	
14	Data classification (public / internal / confidential-unreleased / personal)	
15	Existing-solution check (repository + AI Garage searched; near-matches found)	
16	Desired timeline & reach (studio / region / global)	

Appendix B — Evaluation Scorecard

Criterion (weight)	Score 1	Score 3	Score 5
Strategic alignment (15)	Unrelated to Design/enterprise AI priorities	Supports a priority indirectly	Directly advances a stated priority
Global applicability (15)	Single-team niche	Useful to one region	Every studio would use it
Time savings (12)	< 1 h/week/user	2–5 h/week/user	> 5 h/week/user or removes a bottleneck
Cost savings (10)	None identifiable	Moderate, indirect	Material, direct, recurring
Quality improvement (10)	No effect on outcomes	Fewer errors/reworks likely	Demonstrably better design outcomes
Technical feasibility (10)	Needs unavailable tech/skills	Feasible with help (Talent Pool)	Ready on approved platforms now
User experience (8)	Adds friction	Neutral to mildly better	Users actively want it
Security (8) — gating	Unresolvable exposure (score 1 = pause)	Manageable with controls	No new exposure
Data (6) — gating	Data unavailable/unusable	Available with preparation	Available, clean, rights

	(score 1 = pause)		cleared
Change management (6)	Major retraining/disruption	Moderate training need	Near-zero adoption burden

Total = Σ(score × weight) / 5, max 100. Guidance: ≥ 70 proceed · 50–69 park with revisit date · < 50 decline/merge/redirect. Council may override with recorded rationale.

Appendix C — Pilot Charter Template

Section	Content
Initiative & owners	Name · Pilot Owner · Business Sponsor · Process Owner · Region · Tier
Objectives & hypothesis	What we expect to improve and why
Success criteria	Quantified targets with baseline values, agreed with sponsor
Duration	Start / end (30–90 days; default 60)
Participants	Named users, minimum group size
Measurement approach	Data collected, method, by whom, how often
Checkpoints	Mid-point Council review date; on-track / at-risk / stop
Security & data notes	Classification, environment, tier-required reviews completed
Exit	Final assessment due date (≤ 2 weeks after end); rollback plan

Appendix D — Final Assessment Template

Section	Content
Results vs success criteria	Each criterion: baseline → target → achieved
Value estimate	Validated hours/€ at pilot scope; projection at scale with assumptions
User feedback	Participant ratings and representative quotes
Security / data observations	Issues encountered; controls that worked
Recommendation	Accelerate · Iterate once (with changes) · Retire
Lessons learned	What the next team should know — mandatory, success or failure
Knowledge assets	Links: charter, data, training drafts, SME contacts

Appendix E — Enterprise Alignment Glossary

Framework term	Enterprise counterpart
Design AI repository	Design Data Product, registered via Collibra / Data&AI Marketplace
Regional AI Champions	AI & Data Ambassadors (Design seats in the 255 → 2,000 network)
Design Data Managers / Knowledge Stewards	"Nominate Data Managers within your domain" (hub-

	leader call-to-action)
Pipeline stages (Explore→Scale)	Compatible with Wave stage gates (L0 Idea → L3 Approved) for VCP-flagged initiatives
Estimated business value	VCP value-creation levers, tracked in Wave where material
Reporting metrics	Enterprise Data&AI product KPIs: usage/adoption · health & quality · incremental value · time to market
Build support for pilots	GDPA Global Platform (standardized foundations) + AI & Data Talent Pool (capacity on request)